


When RMSE Lies: Uncertainty Calibration Reveals GP-Direct Overconfidence at $N = 2$

(a) Calibration Coverage

Coverage drops catastrophically for GP-Direct at $N = 2$



(b) Predictive Sharpness (NLL)

GP-Direct is 3.2× worse at $N = 2$ despite similar RMSE

