



Uncertainty Calibration: Selective MT-GP vs GP-Direct

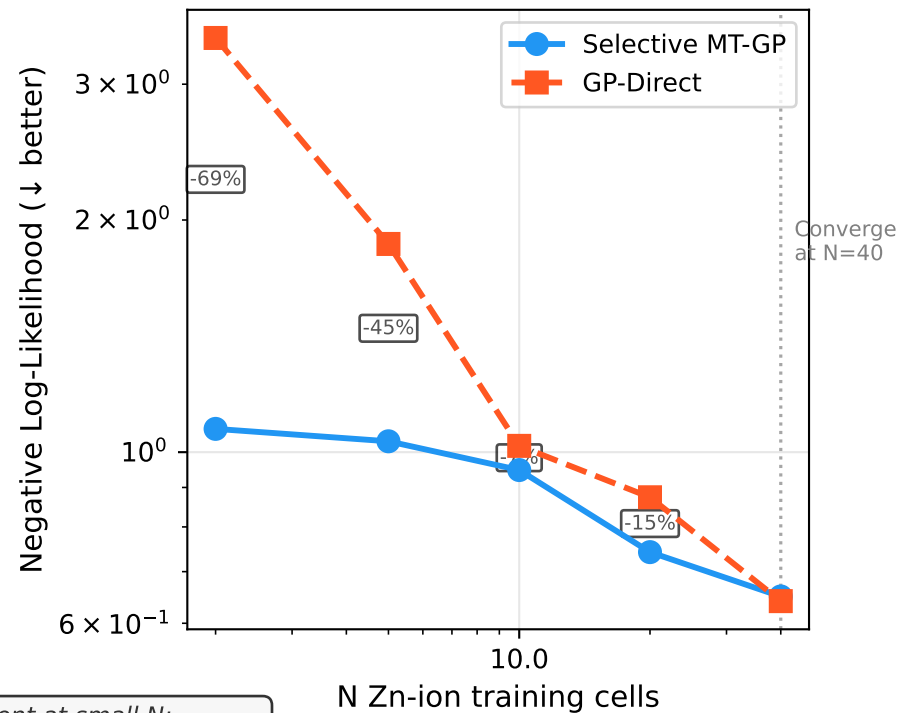
(A) Predictive Accuracy



(B) Calibration: Coverage



(C) Calibration: NLL



GP-Direct achieves lower RMSE (Panel A) but is severely overconfident at small N: 90% coverage drops to 50% at N=2 (Panel B); NLL is $3.2\times$ worse than Selective MT-GP (Panel C). Transfer learning provides engineering safety via calibrated uncertainty, not raw accuracy.